



```
In [11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import sys

# Ensure images are saved to the correct directory
FIGURES_DIR = '../reports/figures'
os.makedirs(FIGURES_DIR, exist_ok=True)

from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
```

1. Helper Functions

Functions to generate standard EDA plots: Feature Importance, Histograms, and Boxplots.

```
In [12]: def save_and_show_feature_importance(df, target, title, filename, top_n=15):
    print(f"Processing Feature Importance: {title}...")
    X = df.drop(target, axis=1)
    y = df[target]

    numeric_features = X.select_dtypes(include=['int64', 'float64']).columns
    categorical_features = X.select_dtypes(include=['object']).columns

    numeric_transformer = SimpleImputer(strategy='median')
    categorical_transformer = Pipeline(steps=[
        ('imputer', SimpleImputer(strategy='most_frequent')),
        ('encoder', OneHotEncoder(handle_unknown='ignore'))
    ])

    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numeric_transformer, numeric_features),
            ('cat', categorical_transformer, categorical_features)
        ]
    )

    clf = RandomForestClassifier(n_estimators=100, random_state=42, n_jobs=-1)
    pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                               ('classifier', clf)])

    pipeline.fit(X, y)

    try:
```

```

        onehot_cols = pipeline.named_steps['preprocessor'].named_transformers_
        feature_names = np.r_[numeric_features, onehot_cols]
    except:
        feature_names = numeric_features

    importances = pipeline.named_steps['classifier'].feature_importances_
    indices = np.argsort(importances)[::-1]

    plt.figure(figsize=(12, 8))
    sns.barplot(x=importances[indices[:top_n]], y=feature_names[indices[:top_n]])
    plt.title(f'Feature Importance: {title}')
    plt.xlabel('Relative Importance')
    plt.tight_layout()
    plt.savefig(os.path.join(FIGURES_DIR, filename))
    plt.show()

def plot_distributions(df, cols, title_prefix, filename_suffix):
    print(f"Plotting Distributions: {title_prefix}...")
    num_cols = len(cols)
    rows = (num_cols // 3) + 1
    plt.figure(figsize=(15, rows * 4))
    for i, col in enumerate(cols):
        plt.subplot(rows, 3, i + 1)
        sns.histplot(df[col], kde=True, bins=30, color='skyblue')
        plt.title(f'Distribution of {col}')
    plt.suptitle(f'{title_prefix} Feature Distributions', y=1.02)
    plt.tight_layout()
    plt.savefig(os.path.join(FIGURES_DIR, f'{filename_suffix}_distributions.png'))
    plt.show()

def plot_boxplots(df, cols, title_prefix, filename_suffix):
    print(f"Plotting Outliers (Boxplots): {title_prefix}...")
    num_cols = len(cols)
    rows = (num_cols // 3) + 1
    plt.figure(figsize=(15, rows * 4))
    for i, col in enumerate(cols):
        plt.subplot(rows, 3, i + 1)
        sns.boxplot(x=df[col], color='salmon')
        plt.title(f'Boxplot of {col}')
    plt.suptitle(f'{title_prefix} Feature Outliers', y=1.02)
    plt.tight_layout()
    plt.savefig(os.path.join(FIGURES_DIR, f'{filename_suffix}_boxplots.png'))
    plt.show()

```

2. Smoker Status EDA

```
In [13]: df_smoker = pd.read_csv('../data/raw/train_dataset.csv')
print("Shape:", df_smoker.shape)
display(df_smoker.describe())
```

Shape: (38984, 23)

	age	height(cm)	weight(kg)	waist(cm)	eyesight(left)	eyesight(right)
count	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000	38984.000000
mean	44.127591	164.689488	65.938718	82.062115	1.014955	1.014955
std	12.063564	9.187507	12.896581	9.326798	0.498527	0.498527
min	20.000000	130.000000	30.000000	51.000000	0.100000	0.100000
25%	40.000000	160.000000	55.000000	76.000000	0.800000	0.800000
50%	40.000000	165.000000	65.000000	82.000000	1.000000	1.000000
75%	55.000000	170.000000	75.000000	88.000000	1.200000	1.200000
max	85.000000	190.000000	135.000000	129.000000	9.900000	9.900000

8 rows × 23 columns

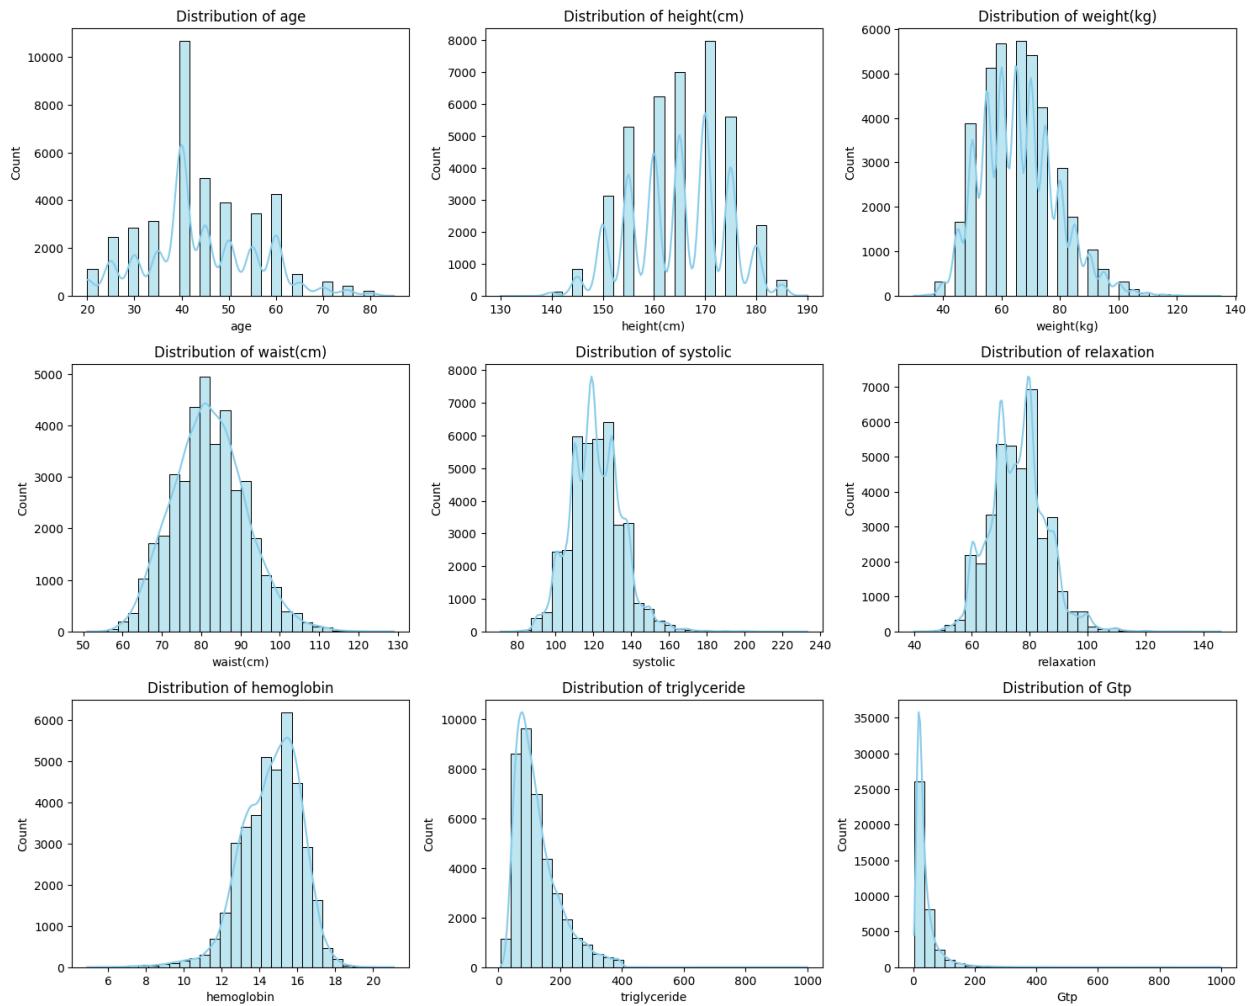
2.1 Feature Distributions (Normality Check)

Crucial for models like SVM/KNN to see if data is skewed.

```
In [14]: numeric_cols_smoker = df_smoker.select_dtypes(include=['float64', 'int64']).columns
# Select top 9 interesting physiological features to avoid clutter
selected_cols = ['age', 'height(cm)', 'weight(kg)', 'waist(cm)', 'systolic', 'diastolic', 'smoker', 'Smoker', 'smoker']
plot_distributions(df_smoker, selected_cols, 'Smoker', 'smoker')
```

Plotting Distributions: Smoker...

Smoker Feature Distributions

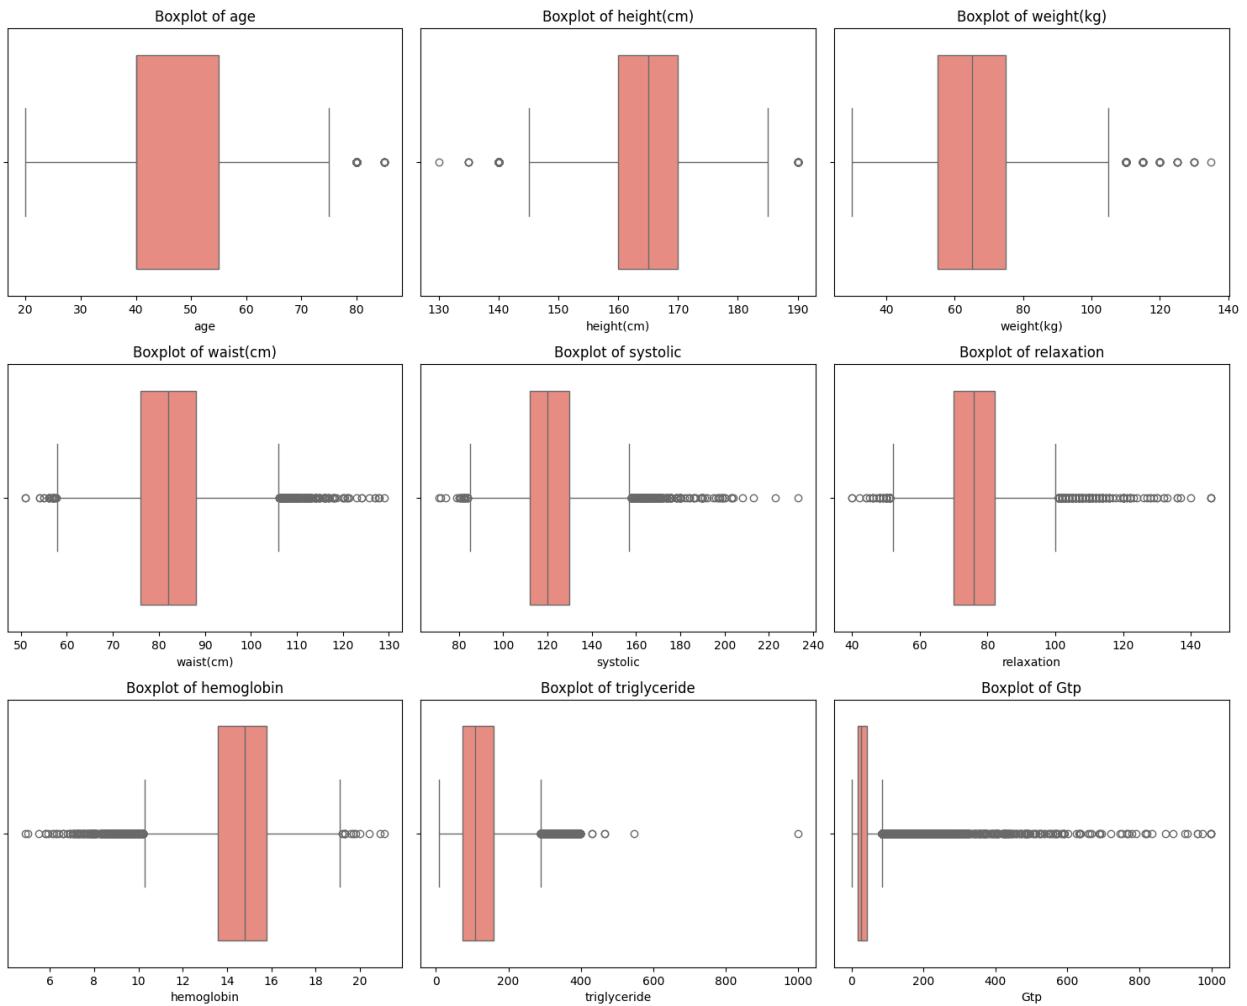


2.2 Outlier Detection

```
In [15]: plot_boxplots(df_smoker, selected_cols, 'Smoker', 'smoker')
```

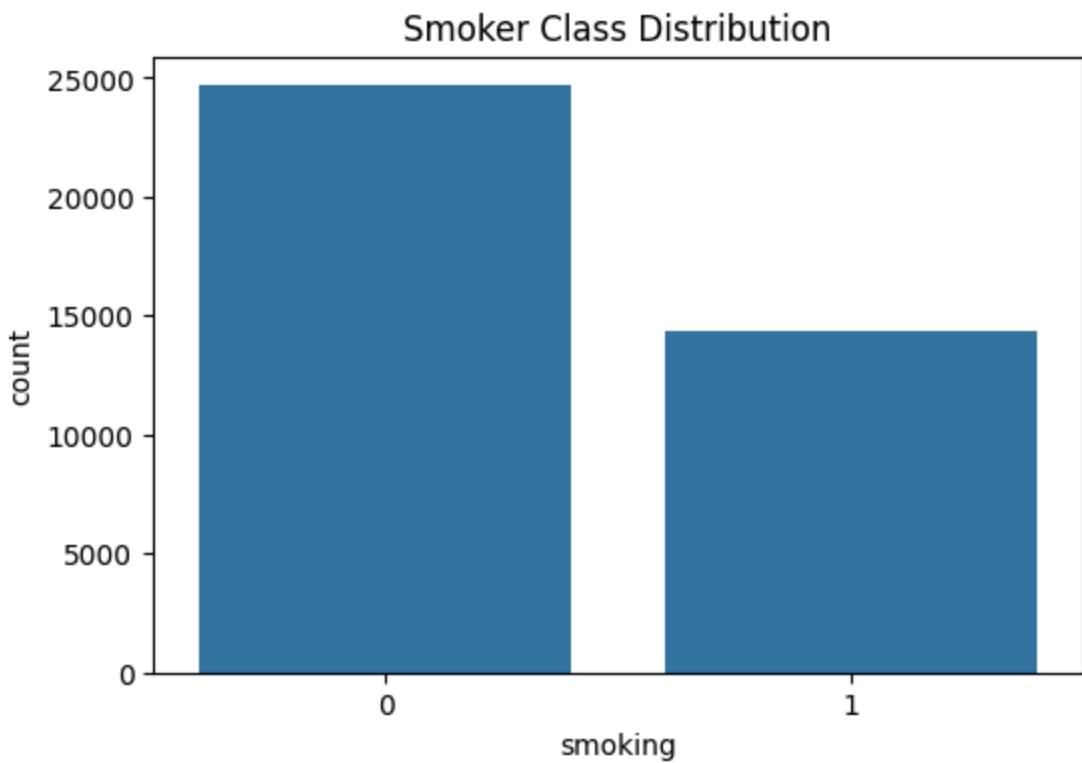
Plotting Outliers (Boxplots): Smoker...

Smoker Feature Outliers



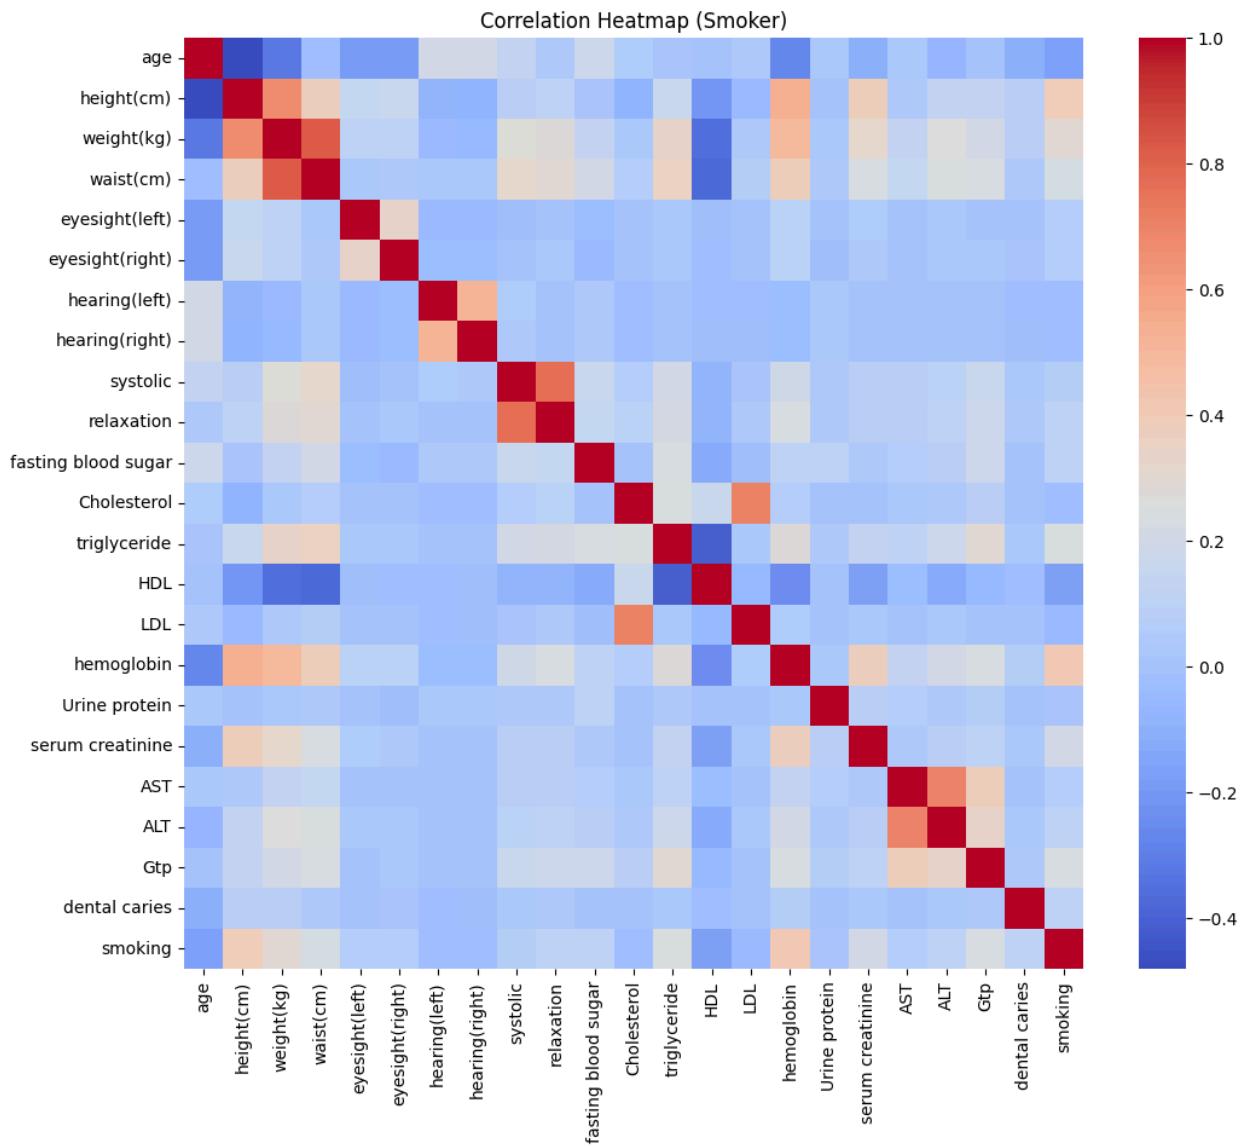
2.3 Class Balance

```
In [16]: plt.figure(figsize=(6, 4))
sns.countplot(x='smoking', data=df_smoker)
plt.title('Smoker Class Distribution')
plt.savefig(os.path.join(FIGURES_DIR, 'smoker_class_dist.png'))
plt.show()
```



2.4 Correlation Heatmap

```
In [17]: plt.figure(figsize=(12, 10))
sns.heatmap(df_smoker.corr(numeric_only=True), cmap='coolwarm', annot=False)
plt.title('Correlation Heatmap (Smoker)')
plt.savefig(os.path.join(FIGURES_DIR, 'smoker_correlation.png'))
plt.show()
```



2.5 Feature Importance

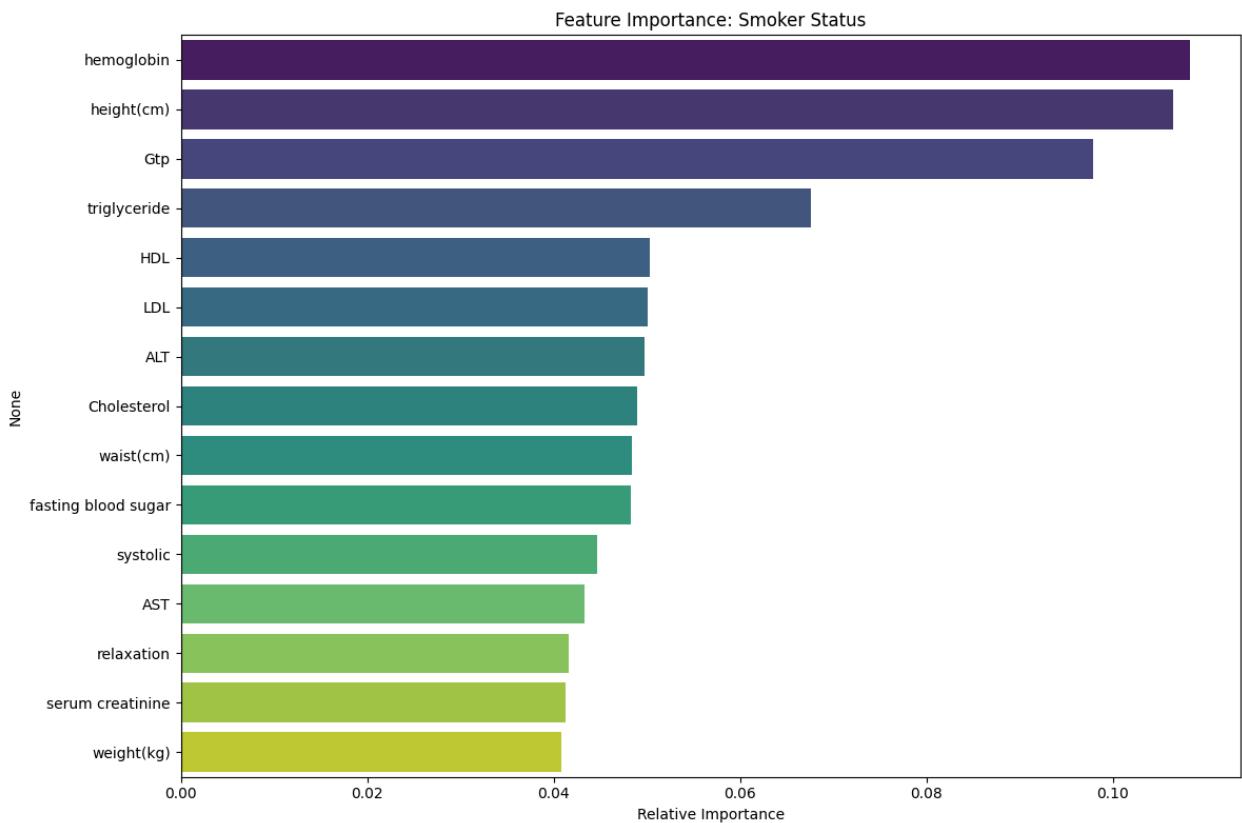
```
In [18]: save_and_show_feature_importance(df_smoker, 'smoking', 'Smoker Status', 'featu
```

Processing Feature Importance: Smoker Status...

C:\Users\NIKHIL AGRAWAL\AppData\Local\Temp\ipykernel_6744\1093813786.py:37: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=importances[indices[:top_n]], y=feature_names[indices[:top_n]], palette='viridis')
```



3. Forest Cover EDA

```
In [19]: df_forest = pd.read_csv('../data/raw/covtype.csv')
if 'Id' in df_forest.columns:
    df_forest = df_forest.drop('Id', axis=1)
print("Shape:", df_forest.shape)
display(df_forest.describe())
```

Shape: (581012, 55)

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydr
count	581012.000000	581012.000000	581012.000000	581012.00
mean	2959.365301	155.656807	14.103704	269.42
std	279.984734	111.913721	7.488242	212.54
min	1859.000000	0.000000	0.000000	0.00
25%	2809.000000	58.000000	9.000000	108.00
50%	2996.000000	127.000000	13.000000	218.00
75%	3163.000000	260.000000	18.000000	384.00
max	3858.000000	360.000000	66.000000	1397.00

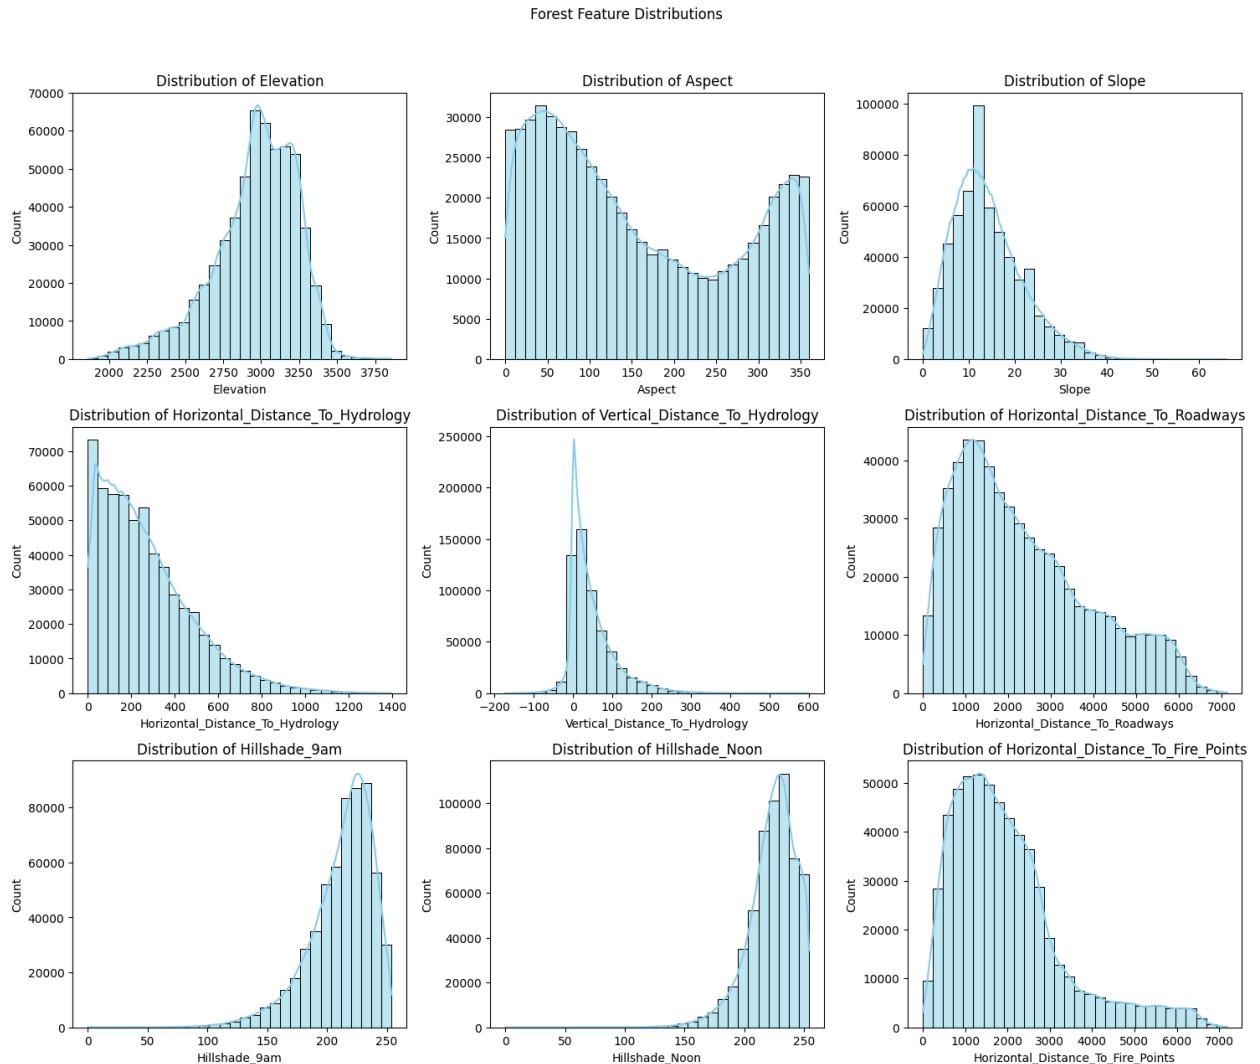
8 rows × 55 columns

3.1 Feature Distributions (Continuous)

Forest cover has many binary columns (Soil types). We focus plots on continuous geographic features.

```
In [20]: geo_cols = ['Elevation', 'Aspect', 'Slope', 'Horizontal_Distance_To_Hydrology',
                 'Vertical_Distance_To_Hydrology', 'Horizontal_Distance_To_Roadways',
                 'Hillshade_9am', 'Hillshade_Noon', 'Horizontal_Distance_To_Fire_Points']
plot_distributions(df_forest, geo_cols, 'Forest', 'forest')
```

Plotting Distributions: Forest...

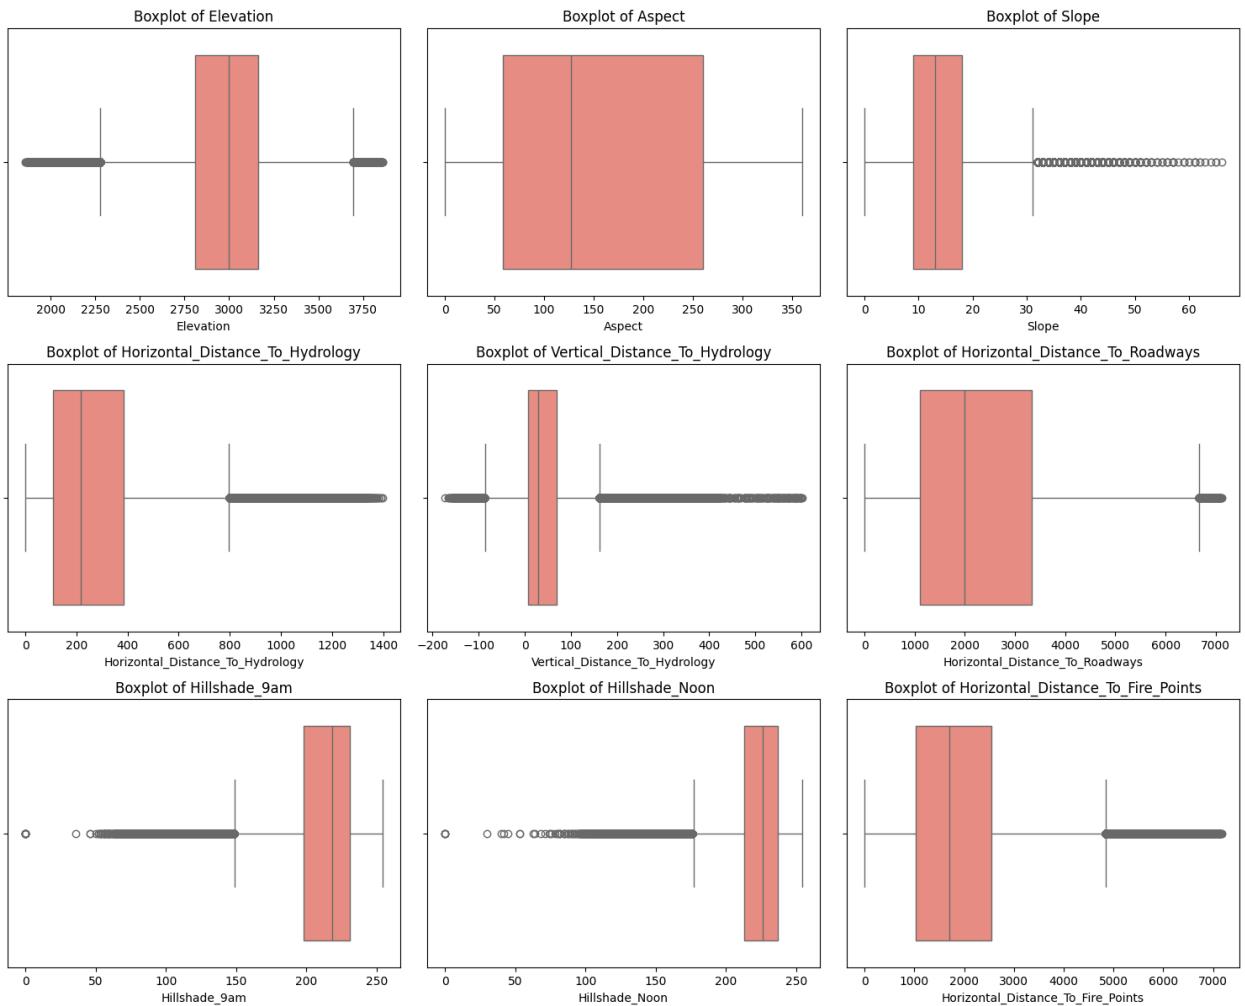


3.2 Outlier Detection

```
In [21]: plot_boxplots(df_forest, geo_cols, 'Forest', 'forest')
```

Plotting Outliers (Boxplots): Forest...

Forest Feature Outliers



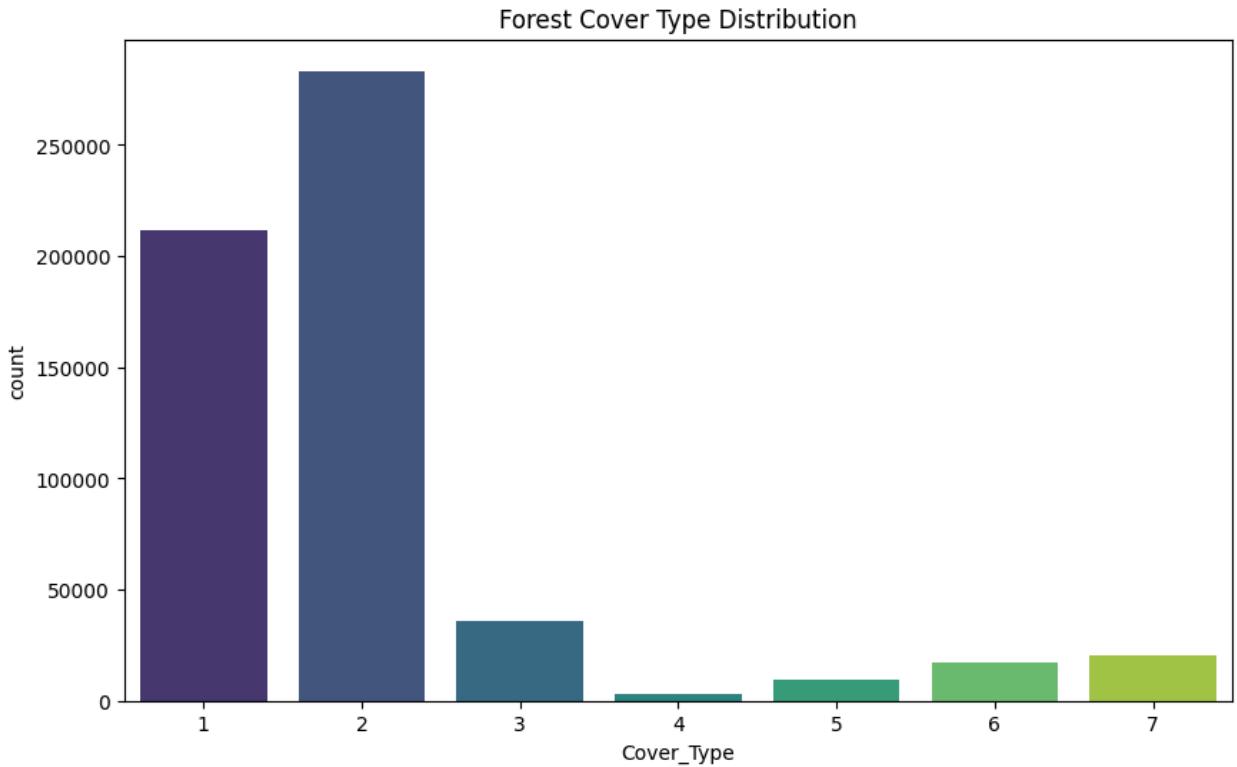
3.3 Class Balance

```
In [22]: plt.figure(figsize=(10, 6))
sns.countplot(x='Cover_Type', data=df_forest, palette='viridis')
plt.title('Forest Cover Type Distribution')
plt.savefig(os.path.join(FIGURES_DIR, 'forest_class_dist.png'))
plt.show()
```

C:\Users\NIKHIL AGRAWAL\AppData\Local\Temp\ipykernel_6744\227076788.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Cover_Type', data=df_forest, palette='viridis')
```



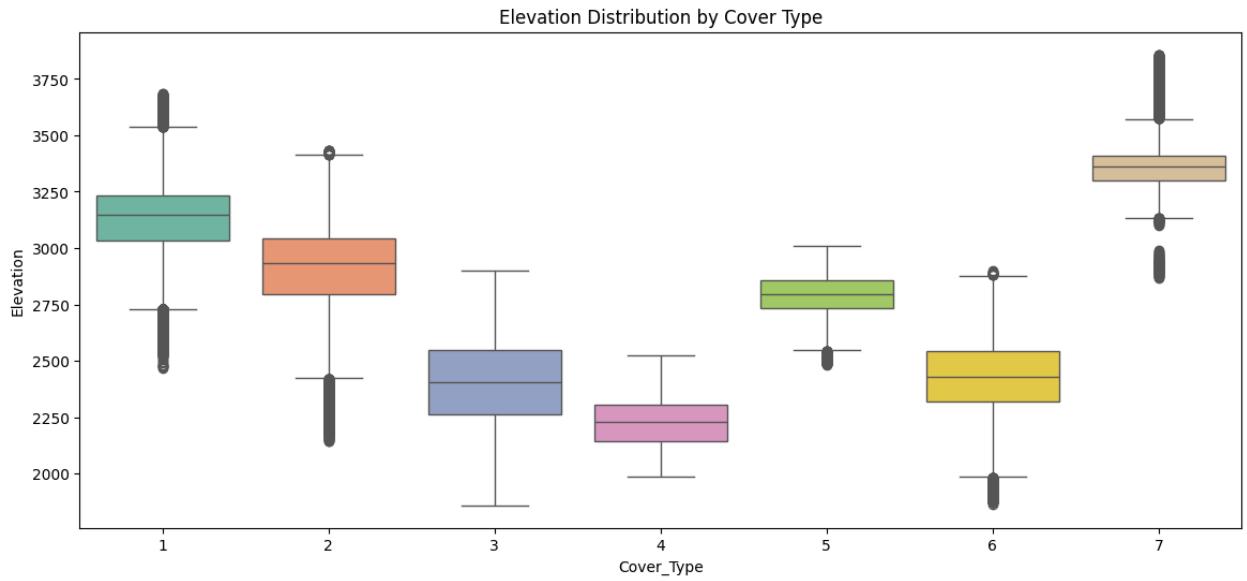
3.4 Key Relationship: Elevation vs Cover Type

```
In [23]: plt.figure(figsize=(14, 6))
sns.boxplot(x='Cover_Type', y='Elevation', data=df_forest, palette='Set2')
plt.title('Elevation Distribution by Cover Type')
plt.savefig(os.path.join(FIGURES_DIR, 'forest_elevation_dist.png'))
plt.show()
```

C:\Users\NIKHIL AGRAWAL\AppData\Local\Temp\ipykernel_6744\620108303.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Cover_Type', y='Elevation', data=df_forest, palette='Set2')
```



3.5 Feature Importance

```
In [24]: save_and_show_feature_importance(df_forest, 'Cover_Type', 'Forest Cover Type',
                                         palette='viridis')

Processing Feature Importance: Forest Cover Type...
C:\Users\NIKHIL AGRAWAL\AppData\Local\Temp\ipykernel_6744\1093813786.py:37: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

sns.barplot(x=importances[indices[:top_n]], y=feature_names[indices[:top_n]],
             palette='viridis')
```

