



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sys
import os
import optuna

sys.path.append(os.path.abspath(os.path.join('..')))

from src.utils import load_data, evaluate_model, compare_models

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import confusion_matrix

optuna.logging.set_verbosity(optuna.logging.WARNING)
```

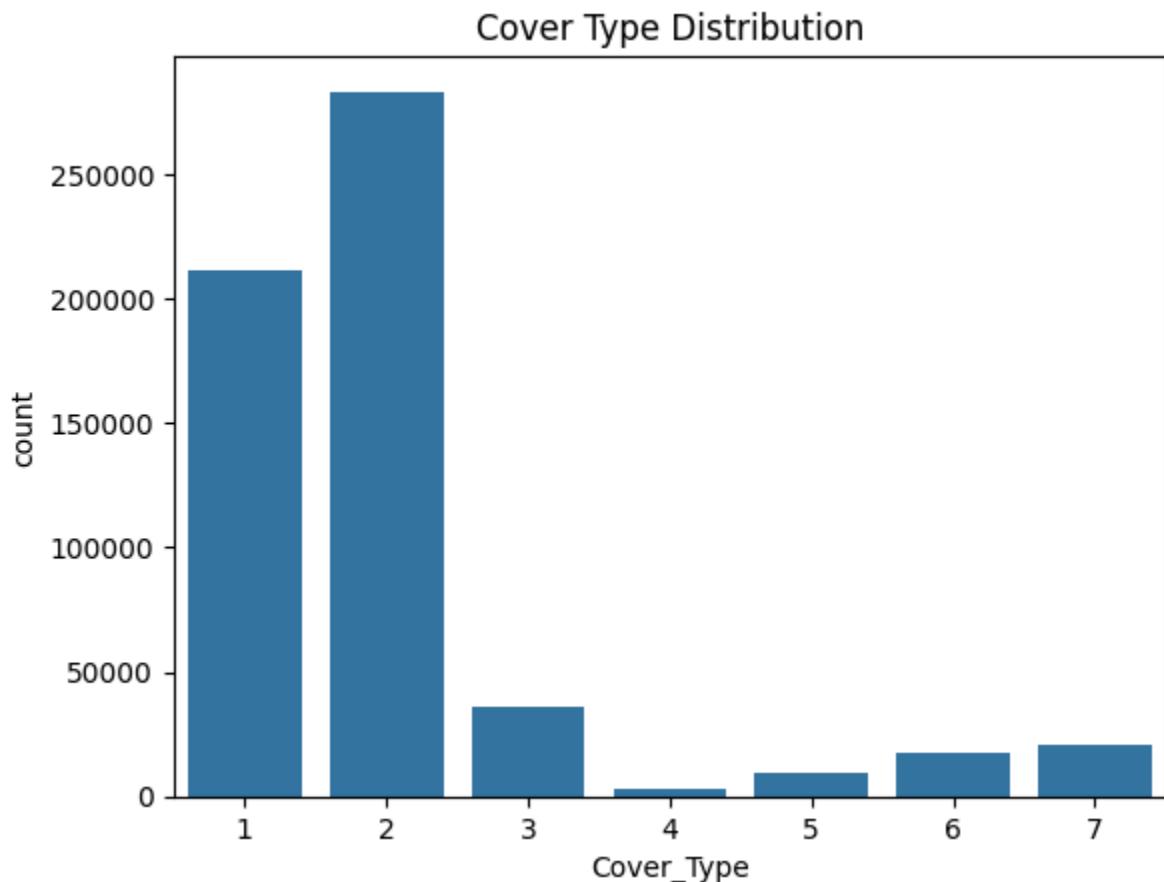
```
C:\Users\NIKHIL AGRAWAL\AppData\Roaming\Python\Python314\site-packages\tqdm\auto.py:21: TqdmWarning: IPython not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```

```
In [2]: df = load_data('../data/raw/covtype.csv')
print(df.info())
target = 'Cover_Type'
```

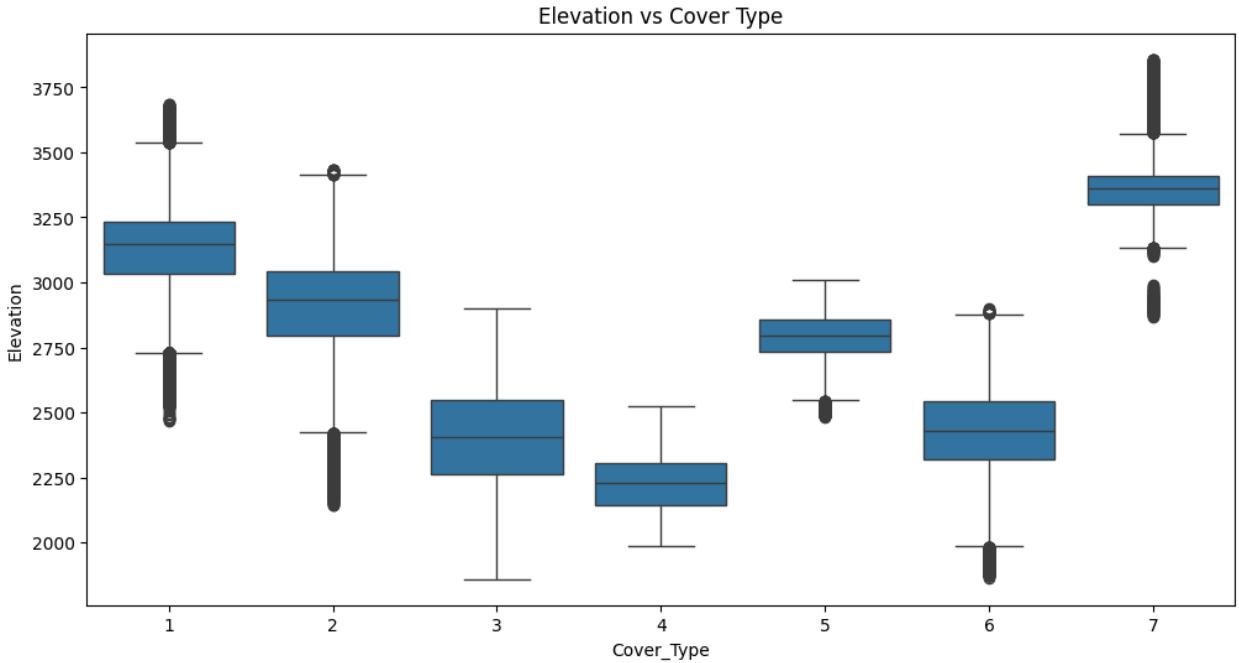
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 581012 entries, 0 to 581011
Data columns (total 55 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   Elevation        581012 non-null int64  
 1   Aspect            581012 non-null int64  
 2   Slope             581012 non-null int64  
 3   Horizontal_Distance_To_Hydrology 581012 non-null int64  
 4   Vertical_Distance_To_Hydrology   581012 non-null int64  
 5   Horizontal_Distance_To_Roadways 581012 non-null int64  
 6   Hillshade_9am       581012 non-null int64  
 7   Hillshade_Noon      581012 non-null int64  
 8   Hillshade_3pm       581012 non-null int64  
 9   Horizontal_Distance_To_Fire_Points 581012 non-null int64  
 10  Wilderness_Areal     581012 non-null int64  
 11  Wilderness_Area2      581012 non-null int64  
 12  Wilderness_Area3      581012 non-null int64  
 13  Wilderness_Area4      581012 non-null int64  
 14  Soil_Type1          581012 non-null int64  
 15  Soil_Type2          581012 non-null int64  
 16  Soil_Type3          581012 non-null int64  
 17  Soil_Type4          581012 non-null int64  
 18  Soil_Type5          581012 non-null int64  
 19  Soil_Type6          581012 non-null int64  
 20  Soil_Type7          581012 non-null int64  
 21  Soil_Type8          581012 non-null int64  
 22  Soil_Type9          581012 non-null int64  
 23  Soil_Type10         581012 non-null int64  
 24  Soil_Type11         581012 non-null int64  
 25  Soil_Type12         581012 non-null int64  
 26  Soil_Type13         581012 non-null int64  
 27  Soil_Type14         581012 non-null int64  
 28  Soil_Type15         581012 non-null int64  
 29  Soil_Type16         581012 non-null int64  
 30  Soil_Type17         581012 non-null int64  
 31  Soil_Type18         581012 non-null int64  
 32  Soil_Type19         581012 non-null int64  
 33  Soil_Type20         581012 non-null int64  
 34  Soil_Type21         581012 non-null int64  
 35  Soil_Type22         581012 non-null int64  
 36  Soil_Type23         581012 non-null int64  
 37  Soil_Type24         581012 non-null int64  
 38  Soil_Type25         581012 non-null int64  
 39  Soil_Type26         581012 non-null int64  
 40  Soil_Type27         581012 non-null int64  
 41  Soil_Type28         581012 non-null int64  
 42  Soil_Type29         581012 non-null int64  
 43  Soil_Type30         581012 non-null int64  
 44  Soil_Type31         581012 non-null int64  
 45  Soil_Type32         581012 non-null int64  
 46  Soil_Type33         581012 non-null int64  
 47  Soil_Type34         581012 non-null int64  
 48  Soil_Type35         581012 non-null int64
```

```
49  Soil_Type36          581012 non-null  int64
50  Soil_Type37          581012 non-null  int64
51  Soil_Type38          581012 non-null  int64
52  Soil_Type39          581012 non-null  int64
53  Soil_Type40          581012 non-null  int64
54  Cover_Type           581012 non-null  int64
dtypes: int64(55)
memory usage: 243.8 MB
None
```

```
In [3]: sns.countplot(x=target, data=df)
plt.title('Cover Type Distribution')
plt.show()
```



```
In [4]: plt.figure(figsize=(12, 6))
sns.boxplot(x=target, y='Elevation', data=df)
plt.title('Elevation vs Cover Type')
plt.show()
```



```
In [5]: if 'Id' in df.columns:
    df = df.drop('Id', axis=1)

X = df.drop(target, axis=1)
y = df[target]

numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns

preprocessor = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

print("Preprocessing data...")
X_processed = preprocessor.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_processed, y, test_size=
```

Preprocessing data...

```
In [6]: # Universal Subsampling for Tuning Phase
# We use a fixed subset of 20,000 samples for all hyperparameter searches to ensure consistency
tuning_size = 20000
idx_tune = np.random.choice(len(X_train), size=tuning_size, replace=False)
X_tune = X_train[idx_tune]
y_tune = y_train.iloc[idx_tune]

print(f"Tuning on subset of {tuning_size} samples for speed.")
```

Tuning on subset of 20000 samples for speed.

```
In [7]: # === 1. SVM Optimization ===
# def objective_svm(trial):
#     C = trial.suggest_float('C', 1e-3, 1e2, log=True)
```

```

#     # Limiting to RBF as it's generally best for non-linear tabular data
#     clf = SVC(C=C, kernel='rbf', gamma='scale', random_state=42)
#     score = cross_val_score(clf, X_tune, y_tune, cv=3, scoring='f1_weighted')
#     return score
#
# print("Tuning SVM...")
# study_svm = optuna.create_study(direction='maximize', sampler=optuna.samplers.NSGAII())
# study_svm.optimize(objective_svm, n_trials=10) # Low trials due to high cost
# print("Best SVM Params:", study_svm.best_params)
# best_svm = SVC(**study_svm.best_params, kernel='rbf', random_state=42)

# Reverting to untuned model as requested
print("Skipping SVM Tuning (Reverted to Default)")
best_svm = SVC(kernel='rbf', random_state=42)

```

Skipping SVM Tuning (Reverted to Default)

```

In [8]: # === 2. Random Forest Optimization ===
# def objective_rf(trial):
#     n_estimators = trial.suggest_int('n_estimators', 50, 300)
#     max_depth = trial.suggest_int('max_depth', 5, 30)
#     min_samples_split = trial.suggest_int('min_samples_split', 2, 20)
#     min_samples_leaf = trial.suggest_int('min_samples_leaf', 1, 10)
#
#     clf = RandomForestClassifier(
#         n_estimators=n_estimators,
#         max_depth=max_depth,
#         min_samples_split=min_samples_split,
#         min_samples_leaf=min_samples_leaf,
#         random_state=42,
#         n_jobs=-1
#     )
#     score = cross_val_score(clf, X_tune, y_tune, cv=3, scoring='f1_weighted')
#     return score
#
# print("Tuning Random Forest...")
# study_rf = optuna.create_study(direction='maximize', sampler=optuna.samplers.NSGAII())
# study_rf.optimize(objective_rf, n_trials=15)
# print("Best RF Params:", study_rf.best_params)
# best_rf = RandomForestClassifier(**study_rf.best_params, random_state=42, n_estimators=100)

# Reverting to untuned model as requested
print("Skipping RF Tuning (Reverted to Default)")
best_rf = RandomForestClassifier(random_state=42, n_jobs=-1)

```

Skipping RF Tuning (Reverted to Default)

```

In [9]: # === 3. KNN Optimization ===
# def objective_knn(trial):
#     n_neighbors = trial.suggest_int('n_neighbors', 3, 20)
#     weights = trial.suggest_categorical('weights', ['uniform', 'distance'])
#     p = trial.suggest_categorical('p', [1, 2]) # 1=Manhattan, 2=Euclidean
#
#     clf = KNeighborsClassifier(n_neighbors=n_neighbors, weights=weights, p=p)

```

```

#     score = cross_val_score(clf, X_tune, y_tune, cv=3, scoring='f1_weighted')
#     return score
#
# print("Tuning KNN...")
# study_knn = optuna.create_study(direction='maximize', sampler=optuna.sampler)
# study_knn.optimize(objective_knn, n_trials=15)
# print("Best KNN Params:", study_knn.best_params)
# best_knn = KNeighborsClassifier(**study_knn.best_params, n_jobs=-1)

# Reverting to untuned model as requested
print("Skipping KNN Tuning (Reverted to Default)")
best_knn = KNeighborsClassifier(n_jobs=-1)

```

Skipping KNN Tuning (Reverted to Default)

```

In [10]: # Final Training Dictionary
models = {
    'Logistic Regression': LogisticRegression(random_state=42, max_iter=1000),
    'KNN': best_knn,
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Random Forest': best_rf,
    'MLP': MLPClassifier(random_state=42, max_iter=200),
    'SVM': best_svm
}

results = {}
trained_models = {}

print("Starting Final Training...")
train_subset_size = 50000 # Use larger subset for final training to keep it re

for name, model in models.items():
    print(f"Training {name}...")

    # SVM and KNN can be very slow on 400k+ rows. We train on a larger subset
    # RF and Decision Tree can technically handle the full set, but for consistency
    # You can increase train_subset_size if you have time/compute.
    if 'SVM' in name:
        idx = np.random.choice(len(X_train), size=train_subset_size, replace=False)
        model.fit(X_train[idx], y_train.iloc[idx])
    else:
        model.fit(X_train, y_train)

    results[name] = evaluate_model(model, X_test, y_test, is_binary=False)
    trained_models[name] = model

```

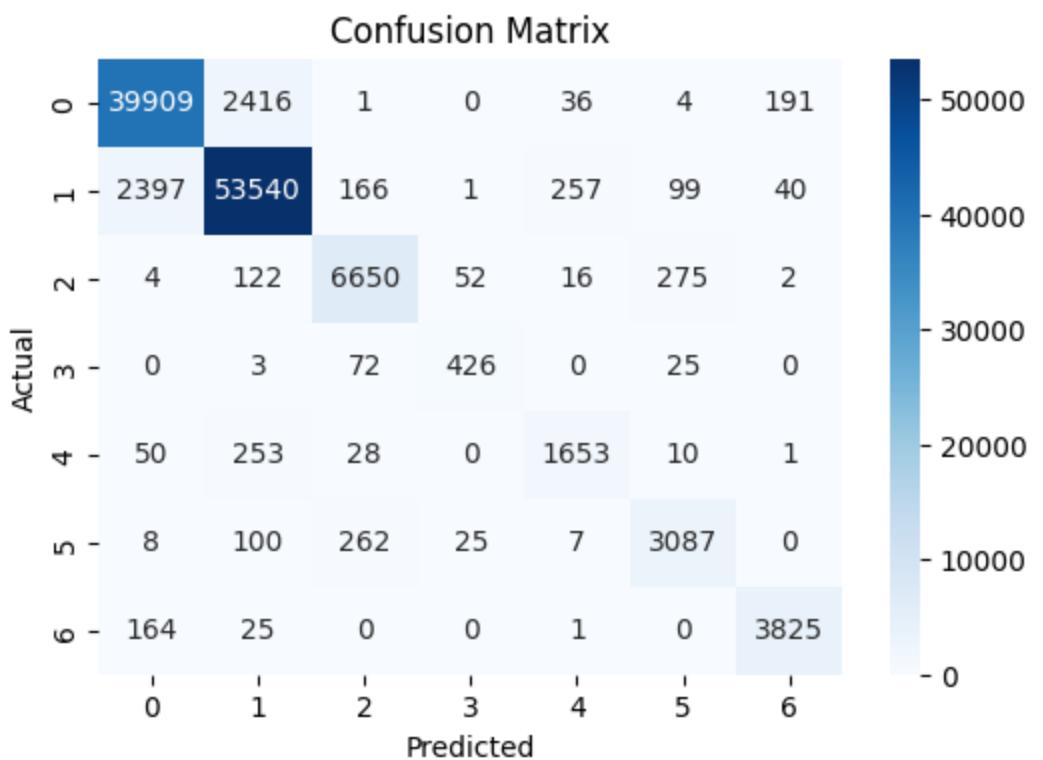
Starting Final Training...  
Training Logistic Regression...

Confusion Matrix							
	0	1	2	3	4	5	6
Actual	29810	11840	8	0	0	17	882
0	10253	45247	633	2	46	291	28
1	0	739	5717	126	4	535	0
2	0	2	255	214	0	55	0
3	12	1890	67	0	16	10	0
4	0	821	1751	29	2	886	0
5	1664	39	0	0	0	0	2312
6	0	1	2	3	4	5	6
Predicted							

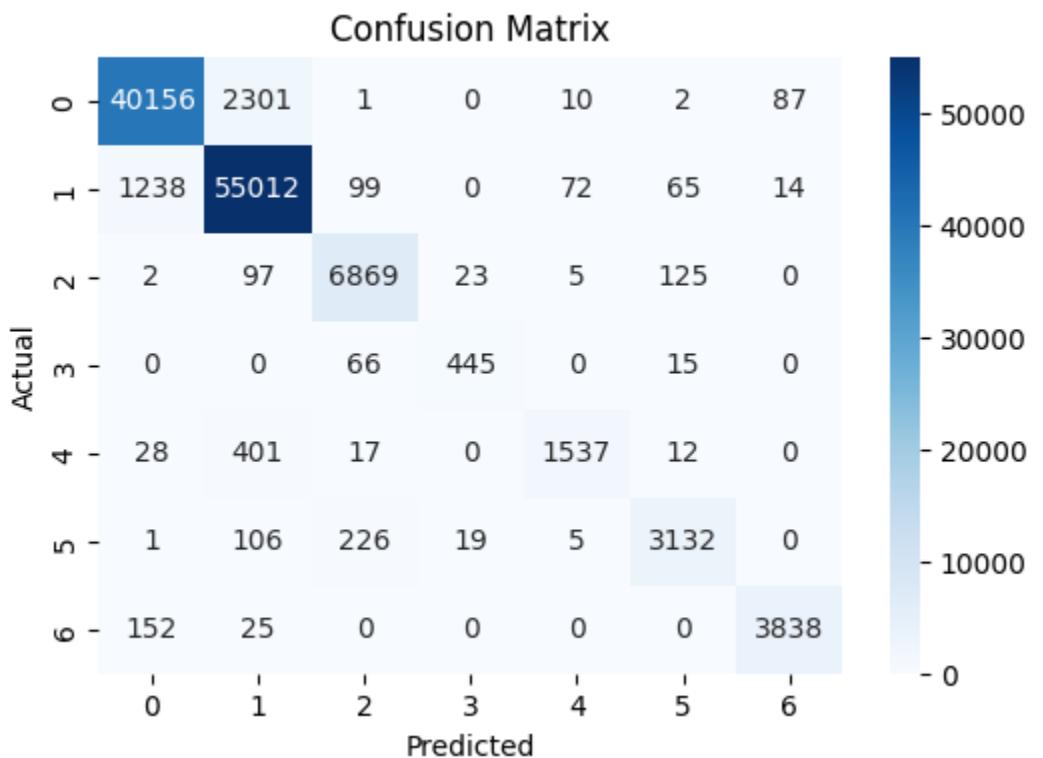
Training KNN...

Confusion Matrix							
	0	1	2	3	4	5	6
Actual	39305	2985	2	0	39	8	218
0	2424	53514	151	0	223	164	24
1	9	190	6525	45	13	339	0
2	0	1	110	382	0	33	0
3	65	368	26	0	1525	11	0
4	14	209	429	25	4	2808	0
5	192	42	0	0	1	0	3780
6	0	1	2	3	4	5	6
Predicted							

Training Decision Tree...

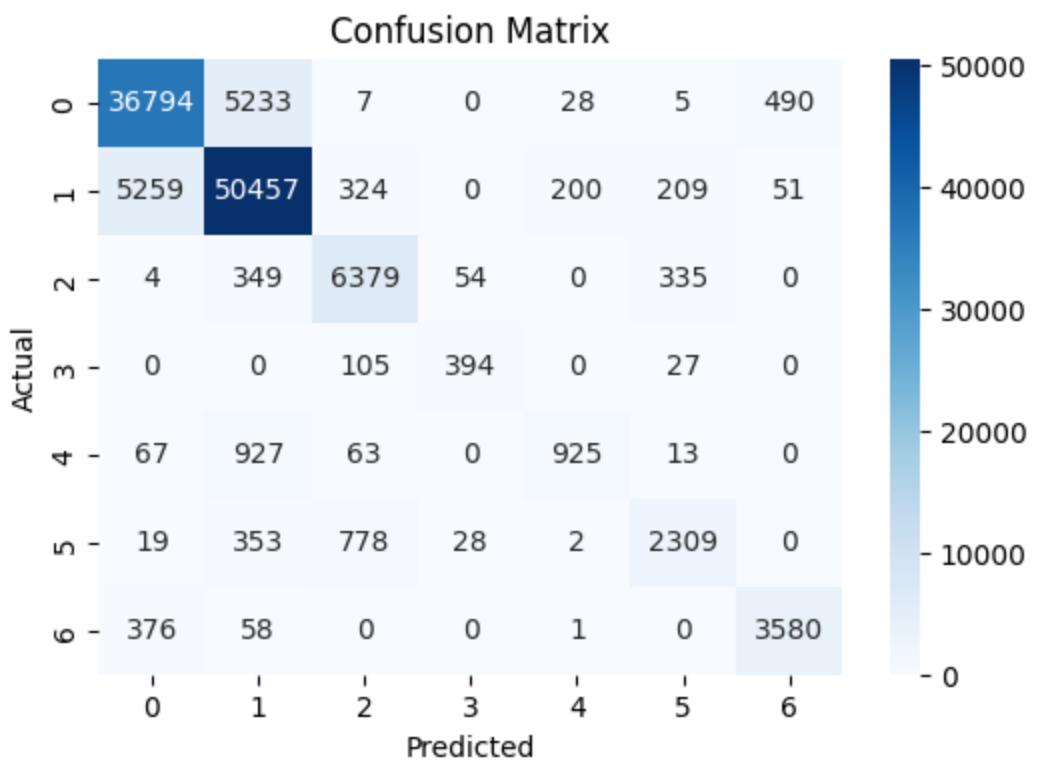


Training Random Forest...

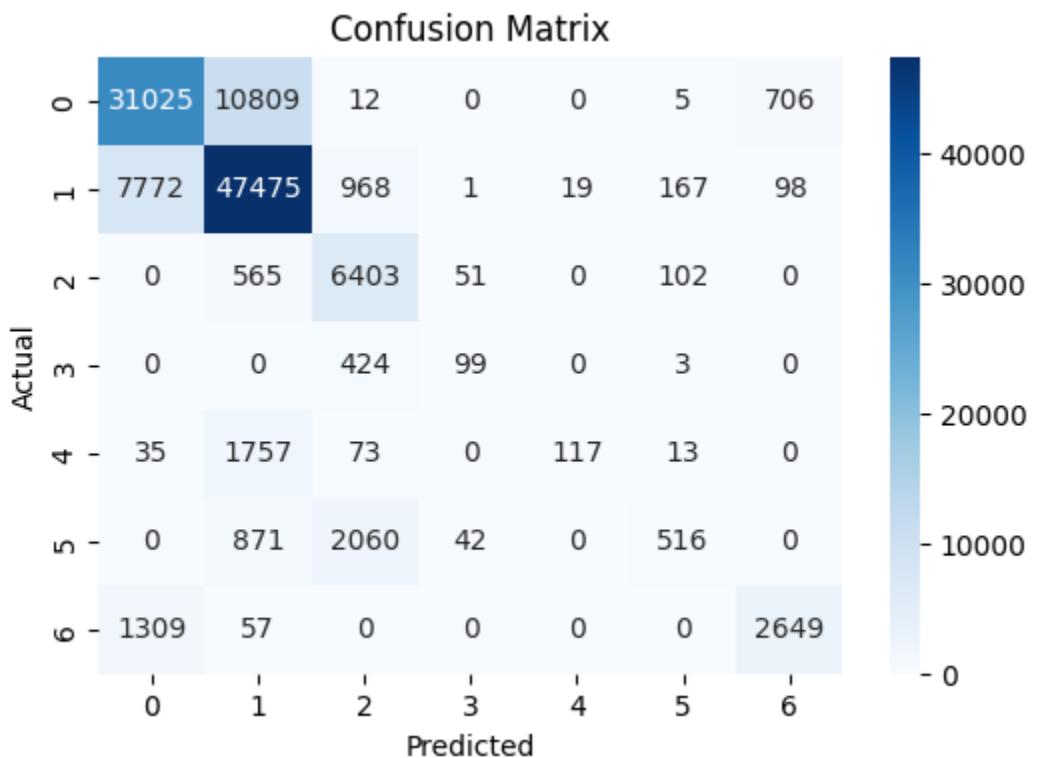


Training MLP...

```
C:\Users\NIKHIL AGRAWAL\AppData\Roaming\Python\Python314\site-packages\sklearn\neural_network\_multilayer_perceptron.py:785: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.
  warnings.warn(
```



Training SVM...



```
In [11]: df_results = compare_models(results)
print(df_results)
```

	Accuracy	Precision	Recall	F1
Random Forest	0.955130	0.955244	0.955130	0.954888
Decision Tree	0.938788	0.938763	0.938788	0.938771
KNN	0.928023	0.927756	0.928023	0.927785
MLP	0.867774	0.866978	0.867774	0.866310
SVM	0.759739	0.760484	0.759739	0.746773
Logistic Regression	0.724611	0.712554	0.724611	0.714467

```
In [12]: best_model_name = df_results.index[0]
best_model = trained_models[best_model_name]

y_pred = best_model.predict(X_test)
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='viridis')
plt.title(f'Multiclass Confusion Matrix ({best_model_name})')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

