



```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sys
import os
import optuna

sys.path.append(os.path.abspath(os.path.join('..')))

from src.utils import load_data, evaluate_model, plot_roc, compare_models

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier


optuna.logging.set_verbosity(optuna.logging.WARNING)
```

```
C:\Users\NIKHIL AGRAWAL\AppData\Roaming\Python\Python314\site-packages\tqdm\auto.py:21: TqdmWarning: IPython not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook_tqdm
```


```
In [2]: df = load_data('../data/raw/train_dataset.csv')
print(df.info())
print(df.duplicated().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38984 entries, 0 to 38983
Data columns (total 23 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   age               38984 non-null   int64  
 1   height(cm)        38984 non-null   int64  
 2   weight(kg)         38984 non-null   int64  
 3   waist(cm)          38984 non-null   float64 
 4   eyesight(left)     38984 non-null   float64 
 5   eyesight(right)    38984 non-null   float64 
 6   hearing(left)       38984 non-null   int64  
 7   hearing(right)      38984 non-null   int64  
 8   systolic           38984 non-null   int64  
 9   relaxation          38984 non-null   int64  
 10  fasting blood sugar 38984 non-null   int64  
 11  Cholesterol        38984 non-null   int64  
 12  triglyceride       38984 non-null   int64  
 13  HDL                38984 non-null   int64  
 14  LDL                38984 non-null   int64  
 15  hemoglobin          38984 non-null   float64 
 16  Urine protein       38984 non-null   int64  
 17  serum creatinine    38984 non-null   float64 
 18  AST                38984 non-null   int64  
 19  ALT                38984 non-null   int64  
 20  Gtp                38984 non-null   int64  
 21  dental caries       38984 non-null   int64  
 22  smoking             38984 non-null   int64  
dtypes: float64(5), int64(18)
memory usage: 6.8 MB
None
5517
```

```
In [3]: target = 'smoking'
sns.countplot(x=target, data=df)
plt.title('Class Balance')
plt.show()
```



```
In [4]: plt.figure(figsize=(12, 10))
sns.heatmap(df.corr(numeric_only=True), cmap='coolwarm', annot=False)
plt.title('Correlation Heatmap')
plt.show()
```



```
In [5]: X = df.drop(target, axis=1)
y = df[target]

categorical_cols = X.select_dtypes(include=['object']).columns
numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns

numerical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
```

```

        ('num', numerical_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)
    )
)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

```

In [6]:

```

# Pre-transform X_train for Optuna efficiency
X_train_transformed = preprocessor.fit_transform(X_train)
X_test_transformed = preprocessor.transform(X_test)

# Use a subset for faster tuning (e.g., 5000 samples)
tuning_size = min(5000, len(X_train_transformed))
idx_tune = np.random.choice(len(X_train_transformed), size=tuning_size, replace=False)
X_tune = X_train_transformed[idx_tune]
y_tune = y_train.iloc[idx_tune]

print(f"Tuning on subset of {tuning_size} samples...")

```

Tuning on subset of 5000 samples...

In [7]:

```

# === Optimization (Commented Out) ===
print("Skipping Tuning (Reverted to Default Models)")

# best_svm = SVC(**study.best_params...)
best_svm = SVC(probability=True, random_state=42)

# best_rf = RandomForestClassifier(**study_rf.best_params...)
best_rf = RandomForestClassifier(random_state=42, n_jobs=-1)

# best_knn = KNeighborsClassifier(**study_knn.best_params...)
best_knn = KNeighborsClassifier(n_jobs=-1)

```

Skipping Tuning (Reverted to Default Models)

In [8]:

```


models = {
    'Logistic Regression': LogisticRegression(random_state=42, max_iter=1000),
    'KNN': best_knn,
    'SVM': best_svm,
    'Naive Bayes': GaussianNB(),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Random Forest': best_rf,
    'MLP': MLPClassifier(random_state=42, max_iter=500)
}

results = {}
trained_models = {}


print("Starting Training...")
for name, model in models.items():
    print(f"Training {name}...")
    model.fit(X_train_transformed, y_train)
    results[name] = evaluate_model(model, X_test_transformed, y_test)
    trained_models[name] = model

```

Starting Training...
Training Logistic Regression...




Training KNN...




Training SVM...

Confusion Matrix




Training Naive Bayes...

Confusion Matrix




Training Decision Tree...

Confusion Matrix




Training Random Forest...

Confusion Matrix



Training MLP...




```
In [9]: df_results = compare_models(results)
print(df_results)
```

	Accuracy	Precision	Recall	F1
Random Forest	0.802873	0.725518	0.732459	0.728972
MLP	0.745158	0.634721	0.697023	0.664415
SVM	0.755034	0.661588	0.661588	0.661588
Decision Tree	0.741183	0.634628	0.671510	0.652548
Naive Bayes	0.697063	0.569319	0.669383	0.615309
KNN	0.718353	0.613078	0.601347	0.607156
Logistic Regression	0.726305	0.634480	0.575124	0.603346

```
In [10]: best_model_name = df_results.index[0]
best_model = trained_models[best_model_name]
print(f"Plotting ROC for Best Model: {best_model_name}")
plot_roc(best_model, X_test_transformed, y_test)
```

Plotting ROC for Best Model: Random Forest



```
In [11]: # === Generate Predictions for Test Dataset ===
print("\n--- Generating Predictions for submission ---")

# 1. Load the test CSV
test_df = load_data('../data/raw/test_dataset.csv')
print(f"Test set loaded: {test_df.shape}")

# 2. Preprocess using the SAME pipeline fitted on train data
X_submission_transformed = preprocessor.transform(test_df)

# 3. Predict using the best model found
# Note: The output is just an array of 0s and 1s
final_predictions = best_model.predict(X_submission_transformed)

# 4. Save to CSV
submission = pd.DataFrame({'prediction': final_predictions})
submission.to_csv('submission_smoker.csv', index=False)
print("Predictions saved to 'submission_smoker.csv'")
print(submission.head())
```

```
--- Generating Predictions for submission ---
Test set loaded: (16708, 22)
Predictions saved to 'submission_smoker.csv'
   prediction
0            1
1            1
2            0
3            0
4            0
```